

gfp_gaussian_process

Likelihood calculation and predictions of 1 dimensional genealogy is re-implementation of the python code:

https://github.com/fioriathos/new_protein_project.

Usage

1 Compile

1.1 Compile locally

The following two libraries are needed:

- nlopt (for minimization)
 - see <https://nlopt.readthedocs.io/en/latest/#download-and-installation>
- Eigen (for linear algebra)
 - see http://eigen.tuxfamily.org/index.php?title=Main_Page

Make sure the correct paths to the two libraries are set in the `Makefile`. Currently both are assumed to be located in the home directory. Then, compile with:

```
cd src; make local
```

1.2 Compile on cluster

1. Install nlopt

- Will install nlopt in home-directory with static linking. You can change that via `DCMAKE_INSTALL_PREFIX`, but make sure to adjust the makefile accordingly!

```
m1 CMake
wget https://github.com/stevengj/nlopt/archive/v2.6.2.tar.gz
tar -xf v2.6.2.tar.gz
cd nlopt-2.6.2/
cmake -DCMAKE_INSTALL_PREFIX=~/.nlopt -DBUILD_SHARED_LIBS=OFF .
make
make install
```

2. Compile

Run `cd src; make cluster`. This will run `m1 GCC/8.3.0`; `m1 Eigen/3.3.7` as well as the compile command! Note, that the modules remain loaded after compilation.

2 Run

```
cd bin
```

`./gfp_gaussian [-options]` with following options:

-h, --help	this help message
-i, --infile	(required) input data file
-b, --parameter_bounds	(required) file setting the type, step, bounds of the parameters
-c, --csv_config	file that sets the cols that will be used from the input file
-l, --print_level	print level >=0, default=0
-o, --outdir	specify output direction and do not use default
-t, --tolerance	absolute tolerance of maximization between optimization steps, default=1e-1
-space, --search_space	search parameter space in 'log' or 'linear' space, default: 'linear'
-stat, --stationary	indicates that the cells are not growing (much)
-beta, --use_beta	indicates that the initial beta will be used to initialize the cells
-m, --maximize	run maximization
-s, --scan	run 1d parameter scan
-p, --predict	run prediction

Example:

```
./gfp_gaussian -c csv_config.txt -b parameter_min.txt -i ../data/simulation_gaussian_gfp.csv -o out/ -l 1 -t 1e-2 -m -p
```

2.1 Required arguments

- `infile` sets the input file that contains the data, eg as given by MOMA (see 2.1.1)
- `parameter_bounds` sets the file that defines the parameter space (see 2.1.2)

2.1.1 Input file

The input file is assumed to fulfill the following:

- the data points of a cell appear as consecutive rows and are in the correct order with respect to time
- the data set has to include all columns that are set via the `csv_config` file, i.e. `time_col`, `length_col`, `fp_col`
- the cells can be uniquely identified via the tags provided via `parent_tags` and `cell_tags` and each mother cell has at most 2 daughter cells. If that is not the case, the `parent_tags` and `cell_tags` are not sufficient and a warning will be printed.
- In order to estimate the initial covariance matrix, the data set needs to contain at least (!) 2 cells.

2.1.2 Parameter file

Syntax for free, bound, fixed (in that order) parameters

- `parameter = init, step`
- `parameter = init, step, lower, upper`
- `parameter = init`

Example:

```
mean_lambda = 0.01, 1e-3
gamma_lambda = 0.01, 1e-3, 1e-4, 0.05
var_lambda = 1e-07
```

ALL parameters are restricted to positive numbers by default avoiding unphysical/meaningless parameter ranges. By setting the lower bound in the parameter file, one can overwrite the lower bounds of the parameters.

The step value is used for the 1d scan to discretize the interval set by lower and upper. During the maximization this will be the initial step size. From nlopt doc:

"For derivative-free local-optimization algorithms, the optimizer must somehow decide on some initial step size to perturb x by when it begins the optimization. This step size should be big enough that the value of the objective changes significantly, but not too big if you want to find the local optimum nearest to x."

2.2 Optional arguments

- `csv_config` sets the file that contains information on which columns will be used from the input file (see 2.2.1)
- `print_level=0` suppresses input of the likelihood calculation, `1` prints every step of the maximization/scan/error bar calculation. This is purely meant for debugging!

- `tolerance` sets the stopping critirium by setting the tolerance of maximization: Stop when an optimization step changes the function value by less than tolerance. By setting very low tolerances one might encounter rounding issues, in that case the last valid step is taken and a warning is printed to stderr.
- `outdir` overwrites default output directory, which is (given the infile `dir/example.csv/`) `dir/example_out/`
- run modes (see 2.2.2)
- `search_space` set the search space of the parameters to be either in log space or linear space. The parameter file does not need to be changed as everything is done internally.
- `stationary` indicates that the cells growth is close to 0. Thus, the initial gfp production is calculated in this limit
- `use_beta` indicates that the bleaching rate beta given in the parameter file as the initial value shall be used to calculate the initial gfp production estimate. Can be used with or without `--stationary` .

2.2.1 Csv_config file

The following settings define how the input file will be read. Default values in brackets.

- `time_col` (`time_sec`): column from which the time is read
- `rescale_time` (60): factor by which time will be divided before the anything is run. This can be used to use a different time unit for the input than for the model parameters.
- `length_col` (`length_um`): column from which the length of the cell is read
- `length_islog` (false): indicates if the cell length in the data file is in logscale (true) or not (false)
- `fp_col` (`gfp_nb`): column from which the intensity is read
- `delm` (,): delimiter between columns, probably ',' or ';'.
- `cell_tag` (`date`, `pos`, `gl`, `id`): columns that will make up the unique cell id, endings like .0 .00 etc of numeric values will be removed
- `parent_tags` (`date`, `pos`, `gl`, `parent_id`): columns that will make up the unique cell id of the parent cell, endings like .0 .00 etc of numeric values will be removed

2.2.2 Run modes

- `m` (`maximize`), `s` (`scan`), `p` (`predict`) will run the respective task. In case `maximize` and `predict` is set, the estimated parameters after the maximization will be used for the prediction. Those parameters that are fixed are of course not effected.
- the 1d parameters scans will calculate the likelihood for the 1d ranges set by the `parameter_bound` file. Note, only "bound" parameters will be scanned.

3 Model parameters

The 2 OU processes are descibed with a mean value (thus the mean growth/production rate), a gamma parameter determining how fast the process is driven towards its mean after a deviation, and a variance that scales the noise term. Thus we have the following parameters including the bleaching rate of the fp, beta:

- Growth rate fluctuations params:
 - `mean_lambda`
 - `gamma_lambda`
 - `var_lambda`
- gfp fluctuation params:
 - `mean_q`
 - `gamma_q`
 - `var_q`
 - `beta`

Additionally, the lenth of the cell and the gfp is effected by measurement noise

- Measurement noise:
 - `var_x`
 - `var_g`

Finally, asymmentric cell division is modeled via two variances of gaussians

- cell division:
 - var_dx
 - var_dg

4 Output

All files that are generated in the maximization and the prediction step are named as follows: `example_f<free>_b<bounds>` and where `<free>` lists the variable via the index as e.g. printed when the code is run and `<bounds>` lists the bound parameters in the same way. Example: `example_f034_b129`. Then, the ending of the different files indicate what they contain.

4.1 Maximization

- Will create 3 files: one for the maximization process (`_iterations.csv`), one for the final estimations (`_final.csv`), and a "parameter file" (`_parameter_file.csv`) that can directly be used as an input for a prediction run
- The iterations file contains the parameter settings at the top 12 lines and all likelihood evaluations of the likelihood maximization
- The final file contains the parameter settings at the top 12 line including a column with the final parameters (i.e. estimated via log likelihood maximization and init value for free parameters) and the estimated error for the estimated parameters via a hessian matrix. The hessian is calculated using a range of finite-differences that are set relative to the value of the respective parameter. I.e. $\epsilon = 1e-2$ corresponds to 1% of each parameter is used for the hessian matrix estimation. Finally, the number of data points, the total log likelihood, the normalized log likelihood, the used optimization algorithm, the set tolerance and the search space (log/linear) is noted.
- The parameter file is in the format of the parameter file that was submitted to the code (see Sec. 2.1.2). It only contains the final estimation of each parameter. Thus, the parameters are treated as "fixed", when the code is run with this parameter file. This file can be used to run the prediction step directly (potentially on a different input file).

4.2 1D scan

- Will create a file for each parameter containing the parameter settings at the top 12 lines and all calculated likelihoods of the scan
- The file (given the input file `example.csv`) is named as follows: `example_<parameter>.csv`, where `<parameter>` is the parameter that is scanned

4.3 Predictions

- Will create 3 files: combined predictions (`_prediction.csv`), backward only (`_prediction_backward.csv`), and forward only (`_prediction_forward.csv`)
- Each file contains the parameter settings at the top 12 lines
- The predictions consist of:
 - the 4 mean quantities of x , g , $1/\lambda$, q
 - the upper triangle of the covariance matrix
 - ... of each time point for each cell in the same order as the input file

Notes

TODO:

- ☒ prepare for cluster
- ☒ write new simulation including asymmetric cell division and tree structure?
- ☐ autocorrelation
- ☒ use log of parameters
- ☒ add stationary option
- ☒ include beta in initializing cells
- ☒ create a parameter file that can be used as input

Technical Notes

Log-Likelihood maximization with multiple cell trees

The log-likelihoods of all cell trees are added and the summed log-likelihood is then maximized. This of course implies that the parameters are the same for all cell trees.

Minimizer

- nlopt
- Note, `total_likelihood` returns `-tl` (negative total log_likelihood). Thus, maximizing the log_likelihood, becomes minimization.
- the wrapper `double total_likelihood(const std::vector<double> ¶ms_vec, std::vector<MOMAdata> &cells);` meant for direct calculation without minimization returns just the log_likelihood

```
double minimize_wrapper(double (*target_func)(const std::vector<double> &x, std::vector<double> &grad, void *p),
                        std::vector<MOMAdata> &cells,
                        Parameter_set &params,
                        double tolerance,
                        bool &found_min)
```

- Current minimizer: LN_COBYLA